Avoiding extrapolation bias when using statistical models to make ecological prediction

Paul B. Conn[1,2] and Devin S. Johnson[1], Peter L. Boveng[1]

[1]*National Marine Mammal Laboratory, NOAA, National Marine Fisheries Service, Alaska Fisheries Science Center, 7600 Sand Point Way NE, Seattle, WA 98115 USA*

Abstract. We'll do this later

Key words: Abundance, Extrapolation, Forecasting, General additive models, Generalized linear models, Independent Variable Hull, Leverage, Occupancy, Spatial regression

Introduction

In ecology and conservation, a common goal is to make predictions about an unsampled random variable given a limited sample from the target population. For instance, given a model ($\mathcal{M}$), estimated parameters ($\hat{\boldsymbol{\theta}}$), and a covariate vector $\mathbf{x}_i$, we often desire to predict a new observation y_{new} at i. For instance, we might use a generalized linear model (McCullagh and Nelder, 1989) or one of its extensions to predict species density or occurrence in a new location, or to predict the future trend of a population.

Early in their training, ecologists and statisticians are warned against extrapolating statistical relationships past the range of observed data. This caution is easily interpreted in the context of single-variable linear regression analysis; one should be cautious in using the fitted relationship to make predictions at some new point y_{new} whenever $x_{new} < \min(\mathbf{x})$ or $x_{new} > \max(\mathbf{x})$. But what about more complicated situations where there are multiple explanatory variables, or when one uses a spatial regression model to account for the residual spatial autocorrelation that is inevitably present in patchy ecological data

[2]Email: paul.conn@noaa.gov

(Lichstein et al., 2002)? How reliable are spatially- or temporally-explicit predictions in sophisticated models for animal abundance and occurrence?

Statisticians have long struggled with the conditions under which fitted regression models are capable of making robust predictions at new combinations of explanatory variables. The issue is sometimes considered more of a philosophical problem than a statistical one, and has even been likened to soothsaying (Ehrenberg and Bound, 1993). To our mind, the reliability of predictions from statistical models is likely a function of several factors, including (i) the intensity of sampling, (ii) spatial or temporal proximity of the prediction location to locations where there are data, (iii) variability of the ecological process, and (iv) the similarity of explanatory covariates in prediction locations when compared to the ensemble of covariates for observed data locations.

Our aim in this paper is to investigate extrapolation bias in the generalized linear model and its extensions, including generalized additive models (GAMs; Hastie and Tibshirani, 1999; Wood, 2006) and spatial, temporal, or spatio-temporal regression models (STRMs). In particular, we exploit some of the same ideas used in multiple linear regression regarding leverage and outliers (Cook, 1979) to operationally define "extrapolation" as making predictions that occur outside of a generalized independent variable hull (gIVH) of observed data points. Application of the gIVH can provide intuition regarding the reliability of predictions in unobserved locations, and can aid in model construction. Also, since the gIVH can be constructed solely with knowledge of sampled locations and explanatory covariates (i.e., it does not necessarily require any observed response variables), it can also be used to help guide survey design. We illustrate use of the gIVH on a simulated occupancy dataset, on a species distribution model (SDM) for ribbon seals in the eastern Bering Sea, and on a population trend model for Steller Sea

Lions (*Phoca fasciata*).

The generalized independent variable hull (gIVH)

Extrapolation is often distinguished from interpolation. In a prediction context, we might define (admittedly quite imprecisely) that extrapolation consists of making predictions that are "outside the range of observed data" while interpolation consists of making predictions "inside the range of observed data." But what exactly do we mean by "outside the range of observed data"? Predictions outside the range of observed covariates? Predictions for locations that are so far from places where data are gathered that we are skeptical that the estimated statistical relationship still holds? To help guide our choice of an operational definition, we turn to early work on outlier detection in simple linear regression analysis.

In the context of outlier detection, Cook (1979) defined an independent variable hull (IVH) as the smallest convex set containing all design points of a full-rank linear regression model. Linear regression models are often written in matrix form, i.e.

$$\mathbf{Y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon},$$

where $\mathbf{Y}$ give observed data, $\mathbf{X}$ is a so-called design matrix that includes explanatory variables (see e.g. Draper and Smith, 1966), and $\boldsymbol{\epsilon}$ represent normally distributed residuals (here and throughout the paper, bold symbols will be used to denote vectors and matrices). Under this formulation, the IVH is defined relative to the hat matrix, $\mathbf{V}_{\text{LR}} = \mathbf{X}(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'$ (where the subscript "LR" denotes linear regression). Letting v denote the maximum diagonal element of $\mathbf{V}_{LR}$, one can examine whether a new design

point, $\mathbf{x}_0$ is within the IVH. In particular, $\mathbf{x}_0$ is within the IVH whenever

$$\mathbf{x}_0'(\mathbf{X}'\mathbf{X})^{-1}\mathbf{x}_0 \leq v. \tag{1}$$

Cook (1979) used this concept to identify influential observations and possible outliers, arguing that design points near the edge of the IVH are deserving of special attention. Similarly, points outside the IVH should be interpreted with caution.

We simulated two sets of design data to help illustrate application of the IVH (Fig. 2). In simple linear regression with one predictor variable, predictions on a hypothetical response variable obtained at covariate values below the lowest observed value or above the highest observed value are primarily outside of the IVH. We suspect this result conforms to most ecologists intuition about what constitutes "extrapolating past the observed data." However, fitting a quadratic model exhibits more nuance; if there is a large gap between design points, it is entirely possible that intermediate covariate values will also be outside of the IVH and thus more likely to result in problematic predictions. Fitting a model with two covariates and both linear and quadratic effects, the shape of the IVH is somewhat more irregular. These simple examples highlight the sometimes counterintuitive nature of the predictive inference problem, a problem that can only become worse as models with more dimensions are contemplated (including those with temporal or spatial structure). Fortunately, the ideas behind the IVH provide potential way forward.

Cook's (1979) formulation for the IVH is particular to linear regression analysis, which assumes iid normally distributed error. Thus, it is not directly applicable to generalized models, such as those including spatial random effects. Further, the hat matrix is not necessarily well defined for models with more general spatial structure. However, since the

hat matrix is proportional to prediction variance, Cook (1979) notes that design points with maximum prediction variance will be located on the boundary of the IVH. Working on the linear predictor scale, we therefore define a generalized independent variable hull (gIVH) as the set of all points $\mathbf{x}$ (note that $\mathbf{x}$ can include both observed and unobserved design points) such that

$$\text{var}(\boldsymbol{\mu}_x|\mathbf{x}) \leq \max[\text{var}(\boldsymbol{\mu}_X|\mathbf{X})], \tag{2}$$

where $\boldsymbol{\mu}_x$ correspond to predictions at $\mathbf{x}$, and $\mathbf{X}$ and $\boldsymbol{\mu}_X$ denote observed design points and predictions at X, respectively.

Generalizations of the linear model are often written in the form

$$Y_i \sim f_Y(g^{-1}(\mu_i)), \tag{3}$$

where f_Y denotes a probability density or mass function (e.g. Bernoulli, Poisson), g gives a link function, and μ_i is a predictor. For many such generalizations, it is possible to specify the μ_i as

$$\boldsymbol{\mu} = \mathbf{X}_{aug}\boldsymbol{\beta}_{aug} + \boldsymbol{\epsilon}, \tag{4}$$

where the $\boldsymbol{\epsilon}$ represent Gaussian error, $\mathbf{X}_{aug}$ denotes an augmented design matrix, and $\boldsymbol{\beta}_{aug}$ denote an augmented vector of parameters. For instance, in a spatial model, $\boldsymbol{\beta}_{aug}$ might include both fixed effect parameters and spatial random effects in a reduced dimension subspace (see Appendix S1 for examples of how numerous types of models can be written in this form).

When models are specified as in Eq. 4, we can write prediction variance generically as

$$\text{var}(\boldsymbol{\mu}_x|\mathbf{x}) = \mathbf{x}\text{var}(\hat{\boldsymbol{\beta}}_{aug})\mathbf{x}', \tag{5}$$

where it is understood that the exact form of $\mathbf{x}$ and $\text{var}(\hat{\boldsymbol{\beta}}_{aug})$ depends on the model chosen (i.e., GLM, GAM, or STRM; Appendix S1). For GLMs, prediction variance is proportional to the original hat matrix $\mathbf{V}_{LR}$ (at least on the linear predictor scale), and Eq. 1 may be used directly. For other models, we have

$$\text{var}(\hat{\boldsymbol{\beta}}_{aug}) = \begin{bmatrix} \boldsymbol{\Sigma}_\beta & \mathbf{0} \\ \mathbf{0} & \boldsymbol{\Sigma}_\alpha \end{bmatrix}, \tag{6}$$

where we simply replace $\boldsymbol{\beta}_{aug}$ with the augmented vector that includes both regression and spatial or smooth parameters (i.e. $\boldsymbol{\beta}$ and $\boldsymbol{\alpha}$; Appendix S1).

The exact form of $\text{var}(\hat{\boldsymbol{\beta}}_{aug})$ differs depending on the underlying model structure and estimation procedure. In the following treatment, we shall focus on Bayesian analysis; although it is not necessarily needed to fit GLMs and GAMs, it helps make STRMs more computationally tractable and puts all of the different models into a common analysis framework. This approach requires specifying priors for model parameters, and prior parameters may appear in expressions for prediction variance. Judicious choices of priors can at times limit the influence of priors on calculation of the gIVH. For example, if we specify the prior on the vector of regression coefficients to be $[\boldsymbol{\beta}] = \text{MVN}(\mathbf{0}, (\tau_\beta X'X)^{-1})$ where τ_β is a fixed constant, a Bayesian implementation of a GLM still results in the gIVH specified in Eq. 1. The ability to use Eq. 5 directly is quite useful, as we only need to know the explanatory variables to be able to diagnose whether predictions lie in or out of the

gIVH. However, the other models considered here (GAMs and STRMs) typically require estimation of an additional precision parameter. In these cases, one solution is to impose prior distributions and integrate over precision parameters (call these $\boldsymbol{\tau}$) to obtain an expectation for $\text{var}(\hat{\boldsymbol{\beta}}_{aug})$. Using the law of the unconscious statistician, we have

$$\text{E}(\text{var}(\hat{\boldsymbol{\beta}}_{aug})) = \int_{\boldsymbol{\tau}} \text{var}(\hat{\boldsymbol{\beta}}_{aug})[\boldsymbol{\tau}] d\boldsymbol{\tau}. \tag{7}$$

Here, and throughout the paper, we use the bracket notation to indicate a probability density function; e.g., $[\boldsymbol{\tau}]$ denotes the joint prior distribution for precision parameters $\boldsymbol{\tau}$. If observations have already been gathered, one could also consider integrating over the marginal posterior distribution using MCMC:

$$\text{E}(\text{var}(\hat{\boldsymbol{\beta}}_{aug})|\mathbf{Y}) = \int_{\boldsymbol{\tau}} \text{var}(\hat{\boldsymbol{\beta}}_{aug})[\boldsymbol{\tau}|\mathbf{Y}] d\boldsymbol{\tau}. \tag{8}$$

We propose to use the gIVH in much the same manner as Cook (1979). In particular, we use the gIVH to differentiate whether spatial predictions are interpolations (predictive design points lying inside the gIVH) or extrapolations (predictive design points lying outside the gIVH). To our knowledge, this is the first time this definition has been used in the context of spatial prediction, but one we shall show can be quite helpful in diagnosing and mitigating extrapolation bias. The gIVH seems ideally situated to this task as it does not necessarily need to rely on gathered response data. Thus, one can examine whether or not prediction points lie within the IVH without ever collecting response data there.

Examples

Simulation study

We conducted simulation analyses to investigate whether the gIVH was useful in diagnosing prediction biases in occupancy and abundance analyses. In both cases, we abundance on a 30×30 grid, and generated as a function of four habitat covariates in addition to spatial process covariance (Appendix S2). The habitat covariates also

Ribbon seal SDM

As part of an international effort, researchers with the U.S. National Marine Fisheries Service conducted aerial surveys over the eastern Bering Sea in 2012 and 2013. Agency scientists used infrared video to detect seals that were on ice, and simultaneous automated digital photographs provided information on species identity. Here, we use spatially referenced count data from photographed ribbon seals, *Phoca fasciata* (Fig. 1) on a subset of 10 flights flown over the Bering Sea in April 2012. These flights were limited to a one week period so that both sea ice conditions and seal distributions could be assumed to be static.

Our objective with this dataset will be to model seal counts on transects through 25km by 25km grid cells as a function of habitat covariates and possible spatial autocorrelation. Estimates of apparent abundance can then be obtained by summing predictions across grid cells. Figure 3 shows the transects flown and the number of ribbon seals encountered in each cell, and Figure 5 show explanatory covariates gathered to help predict ribbon seal abundance. These data are described in fuller detail by Conn et al. (Accepted), who extend the modeling framework of STRMs to account for incomplete detection and species

misidentification errors (see e.g. Conn et al., Accepted). Since our focus in this paper is on illustrating spatial modeling concepts, we devote our efforts to the comparably easier problem of estimating apparent abundance (i.e., uncorrected for vagaries of the detection process).

Inspection of ribbon seal data (Fig. 3) immediately reveals a potential issue with spatial prediction: abundance of ribbon seals appears to be maximized in the southern and/or southeast quadrant of the surveyed area. Predicting abundance in areas further south and east may thus prove problematic. To illustrate, let Y_i denote the ribbon seal count (Y_i) obtained in sampled grid cell i. Suppose that counts arise according to a log-Gaussian Cox process, such that

$$Y_i \sim \text{Poisson}(\lambda_i) \text{ and} \tag{9}$$
$$\log(\lambda_i) = \log(P_i) + \mathbf{X}_i\boldsymbol{\beta} + \eta_i + \epsilon_i,$$

where P_i gives the proportion of area photographed in grid cell i (recall also that $\mathbf{X}_i$ denotes a vector covariates for cell i, $\boldsymbol{\beta}$ are regression coefficients, η_i represents a spatially autocorrelated random effect, and ϵ_i is normally distributed *iid* error).

We could fit any number of predictive models to these data, but we start with a simple generalized linear model where we ignore the spatial random effect, η_i, and use the full suite of predictor covariates (Fig. 5) to fit Eq. 9 to our data. In particular, we fit a model with linear effects of all predictor variables, and with an additional quadratic term for ice concentration (seal density is often maximized at an intermediate value of ice concentration; see Ver Hoef et al., 2013; Conn et al., Accepted). To enable comparison with more complicated types of models, we formulated a generalized Bayesian strategy for

parameter estimation (see Appendix S1). For simplicity, we generated posterior predictions of ribbon seal abundance across the landscape as

$$N_i \sim \text{Poisson}(A_i \lambda_i), \tag{10}$$

where A_i gives the proportion of suitable habitat in cell i (ribbon seals do not inhabit land masses).

Fitting this model to our data,

Steller sea lion trends

Discussion

A number of authors have explored optimal knot placement in spatial models. In the context of predictive process modeling (where a covariance function is specified over a group of knots; see Banerjee et al., 2008), Finley et al. (2009) and Gelfand et al. (2013) considered near-optimal knot placement by minimizing spatially averaged prediction variance. Gelfand et al. knot selection

Using GAs to select knot placement using

Contrast with posterior loss (e.g., Jay's linex loss function estimator)

Contrast with cross validation

David Miller (the evil one)/Simon Wood stuff on edge effects

Contrast with other approaches- Gelfand et al. Bayesian analysis - intrinsic CAR in SDM Chakraborty et al. '10 - spatial abundance modeling - ordinal Latimer et al. 2009 - spatial predictive process modelling in SPDs

Can't be complacent... still possible to get poor/biased results, e.g. if $\tau_\epsilon \to 0$. Can't resolve pathological problems.

Presence-absence data, other link functions (e.g. probit)

extensions to models w/ secondary observation process, measurement error

much attention has been given to collinearity in multiple linear regression - suggest researchers give as much attention to predictive extrapolation bias in predictive models

For predictions with spatial , our experience is that predictions outside the minimum convex polygon where data are obtained can sometimes be more problematic than predictions within the polygon. Spatial prediction surfaces may have a tendency to bend up or down in these areas, resulting in "edge effects" that can lead to positive prediction bias when a log link function is employed (Ver Hoef and Jansen, 2007).

Fieberg PVA

Literature Cited

Banerjee, S., A. E. Gelfand, A. O. Finley, and H. Sang. 2008. Stationary process approximation for the analysis of large spatial datasets. Journal of the Royal Statistical Society B **70**:825–848.

Conn, P. B., J. M. Ver Hoef, B. T. McClintock, E. E. Moreland, J. M. London, M. F. Cameron, S. P. Dahle, and P. L. Boveng. Accepted. Estimating multi-species abundance using automated detection systems: ice-associated seals in the eastern Bering Sea. Methods in Ecology and Evolution .

Cook, R. D. 1979. Influential observations in linear regression. Journal of the American Statistical Association **74**:169–174.

Draper, N. R., and H. Smith. 1966. Applied Regression Analysis. John Wiley & Sons, New York.

Ehrenberg, A., and J. Bound. 1993. Predicatibility and prediction. Journal of the Royal Statistical Society A **156**:167–206.

Finley, A. O., H. Sang, S. Banerjee, and A. E. Gelfand. 2009. Improving the performance of predictive process modeling for large datasets. Computational Statistics & Data Analysis **53**:2873–2884.

Gelfand, A. E., S. Banerjee, and A. O. Finley, 2013. Spatial design for knot selection in knot-based dimension reduction models. *in* J. Mateu and W. G. Muller, editors. Spatio-temporal Design: Advances in Efficient Data Acquisition. Wiley.

Hastie, T. J., and R. J. Tibshirani. 1999. Generalized Additive Models. Chapman & Hall/CRC, Boca Raton, Florida.

Lichstein, J., T. Simons, S. Shriner, and K. E. Franzreb. 2002. Spatial autocorrelation and autoregressive models in ecology. Ecological Monographs **72**:445–463.

McCullagh, P., and J. A. Nelder. 1989. Generalized Linear Models. Chapman and Hall, New York.

Ver Hoef, J., and J. Jansen. 2007. Space-time zero-inflated count models of harbor seals. Environmetrics **18**:697–712.

Ver Hoef, J. M., M. F. Cameron, P. L. Boveng, J. M. London, and E. E. Moreland. 2013. A hierarchical model for abundance of three ice-associated seal species in the eastern Bering Sea. Statistical Methodology page http://dx.doi.org/10.1016/j.stamet.2013.03.001.

Wood, S. N. 2006. Generalized additive models. Chapman & Hall/CRC, Boca Raton, Florida.

FIGURE 1. A ribbon seal, *Phoca fasciata*; the focus of spatial modeling efforts in this paper.

FIGURE 2. Examples IVHs constructed from simulated data. In panels A and B, the investigator plans to model a hypothetical (unmeasured) response variable using a linear regression model as a function of a single covariate, x, obtained at a number of design points (denoted with an "x"). Using x as a simple linear effect (A), only predictions less than the minimum observed value of x or greater than the maximum value of x are outside the IVH (shaded area), as scaled prediction variance in these areas (solid line) are greater than the maximum scaled prediction variance for observed data (dashed line). Using both linear and quadratic effects of x, some intermediate points are also outside the IVH; predictions at these points should also be viewed with caution. Panels C & D show a more complicated IVH when the investigator wishes to relate an unmeasured response variable to linear and quadratic effects of two covariates, x and y, either without interactions (C) or with interactions (D). Any potential predictions in the shaded area are outside of the IVH and should be viewed with caution.

FIGURE 3. Aerial surveys over the Bering Sea in spring of 2012 (blue lines) overlayed on a tesselated surface composed of 25km by 25km grid cells. Gray indicates land, and colored pixels indicate ribbon seal encounters (yellow: 1-2 seals; orange: 3-4 seals; magenta: 5-9 seals; red: 10-15 seals). On average, photographs covered approximately $2.6km^2$ (0.4%) of each surveyed grid cell.

FIGURE 4. Potential covariates gathered to help explain and predict ribbon seal abundance in the eastern Bering Sea. Covariates include distance from mainland

(`dist_mainland`), distance from 1000m depth contour (`dist_shelf`), average remotely sensed sea ice concentration while surveys were being conducted (`ice_conc`), and distance from the southern sea ice edge (`dist_edge`). All covariates except ice concentration were standardized to have a mean of 1.0 prior to plotting and analysis.

FIGURE 5 Posterior median estimates of ribbon seal apparent abundance across the eastern Bering sea for (A) a generalized linear model (GLM), (B) a generalized additive model (GAM), (C) a GLM with known zero data, and (D) a GAM with known zero data. Highlighted cells indicate those where predictive covariate values are outside of the generalized independent variable hull.

Figures

Fig 1

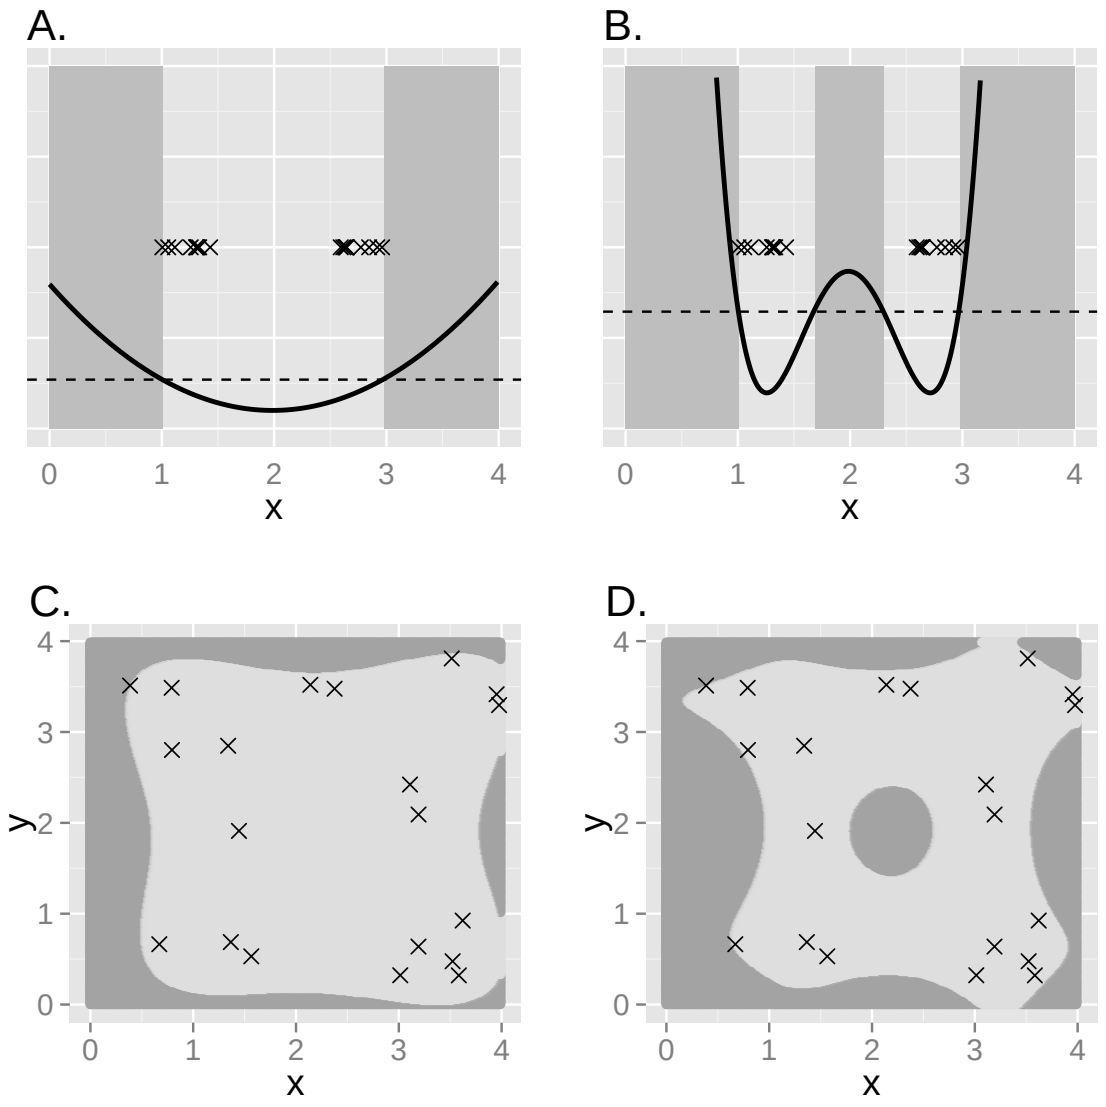

FIG 2

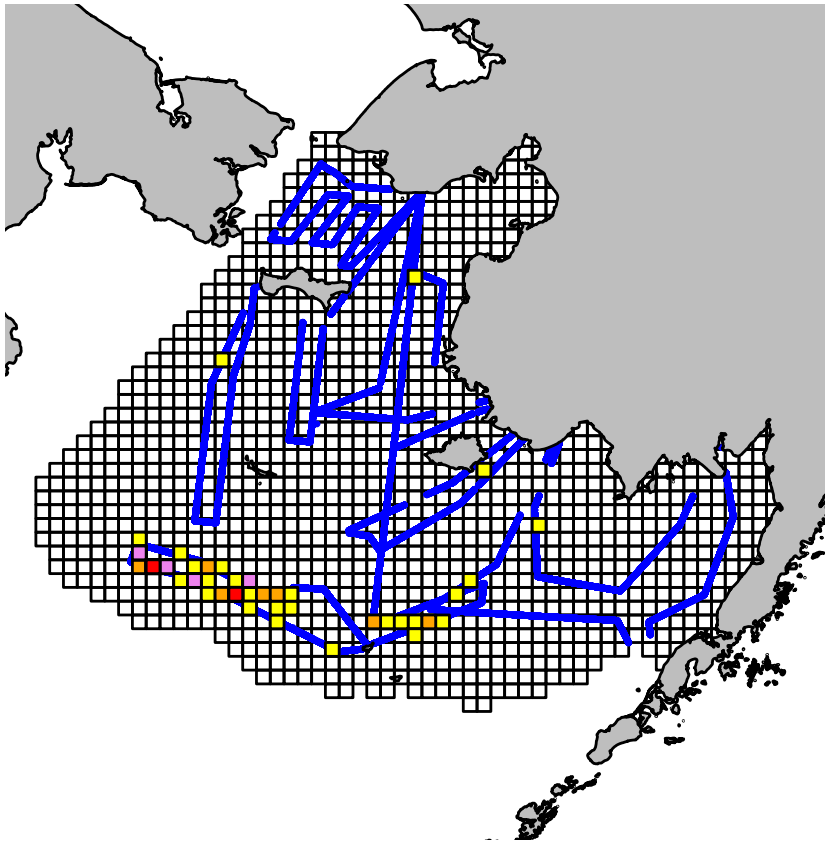

Fig 3

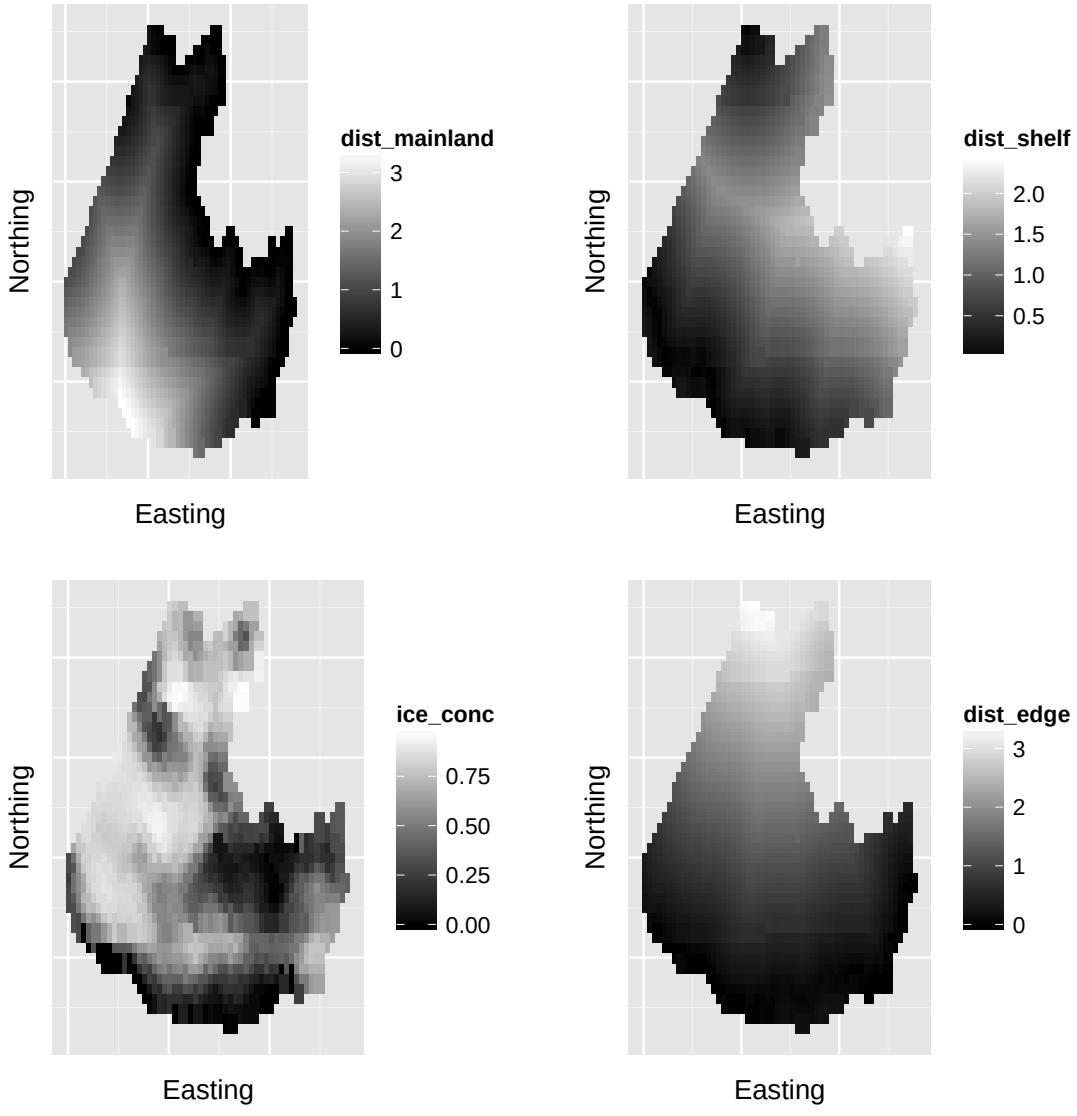
dist_mainland
3
2
1
0
Northing
Easting
dist_shelf
2.0
1.5
1.0
0.5
Northing
Easting
ice_conc
0.75
0.50
0.25
0.00
Northing
Easting
dist_edge
3
2
1
0
Northing
Easting

FIG 4

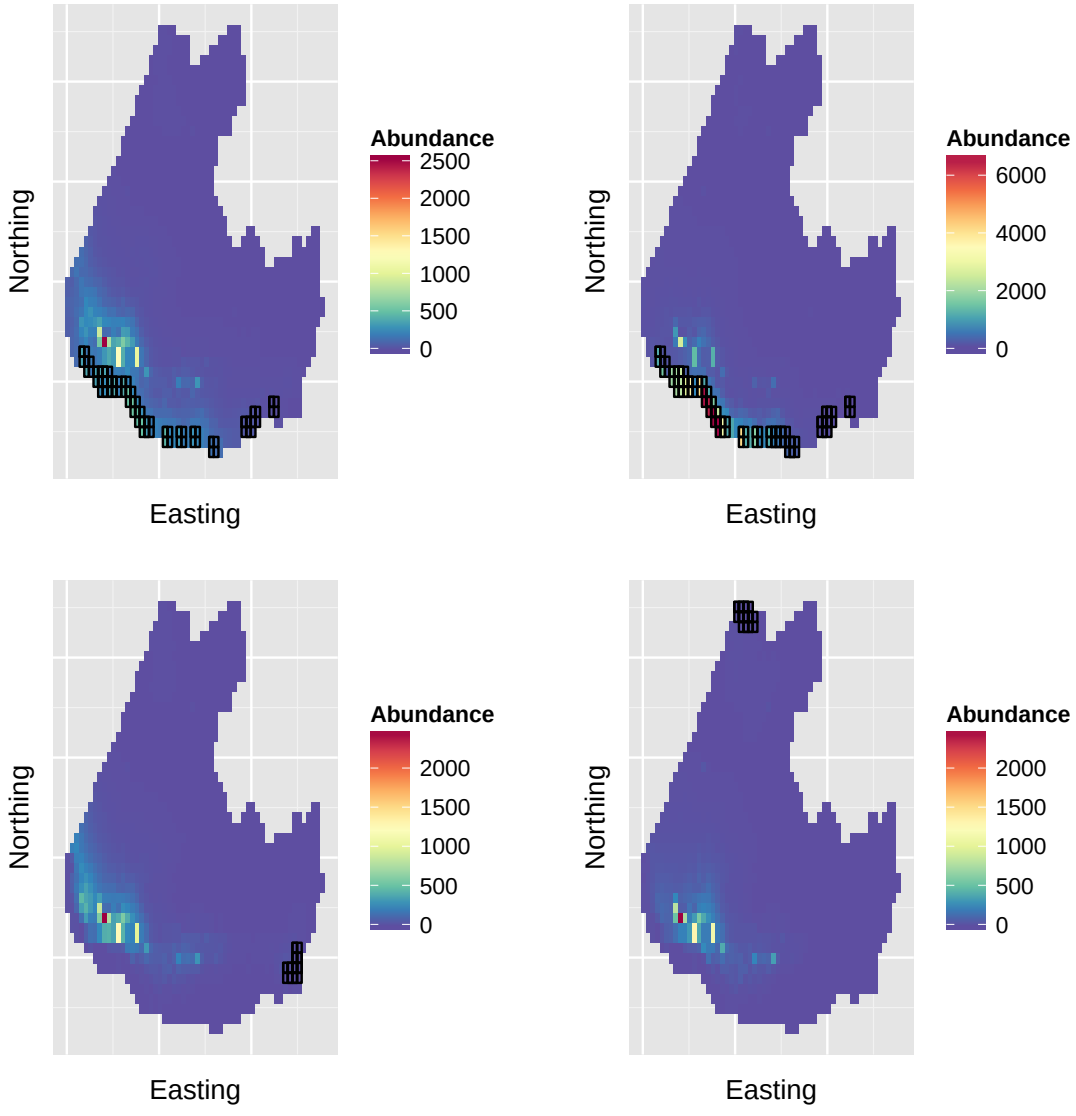
Northing
Easting
Abundance
2500
2000
1500
1000
500
0
Northing
Easting
Abundance
6000
4000
2000
0
Northing
Easting
Abundance
2000
1500
1000
500
0
Northing
Easting
Abundance
2000
1500
1000
500
0

FIG 5